

EC440: Project 2 – Threading Library

Project Goals:

- To understand threads.
- To implement parallel execution within a process

Collaboration Policy:

You are encouraged to discuss this project with your classmates and instructors. You **must** develop your own solution. While you **must not** share your threading code with other students, you **are encouraged** to share test case code that you developed to test your solution for this assignment.

Deadline:

Project 2 is due March 8 at 4:30 PM EDT.

Project Description:

The main deliverable for this project is a basic thread system for Linux. In the lectures, we learned that threads are independent units of execution that run (virtually) in parallel in the address space of a single process. As a result, they share the same heap memory, open files (file descriptors), process identifier, etc. Each thread has its own context, which consists of a set of CPU registers and a stack. The thread subsystem provides a set of library functions that applications may use to create, start and terminate threads, and manipulate them in various ways.

The most well-known and widespread standard that specifies a set of interfaces for multi-threaded programming on Unix-style operating systems is called POSIX threads (or *pthreads*). Note that pthreads merely prescribes the interface of the threading functionality. The implementation of that interface can exist in user-space, take advantage of kernel-mode threads (if provided by the operating system), or mix the two approaches. In this project, you will implement a small subset of the pthread API exclusively in user-mode. In particular, we aim to implement the following three functions from the pthread interface in user mode on Linux (prototypes and explanations partially taken from the respective man pages):

```
int pthread_create(  
    pthread_t *thread,  
    const pthread_attr_t *attr,  
    void *(*start_routine) (void *),  
    void *arg);
```

The `pthread_create()` function creates a new thread within a process. Upon successful completion, `pthread_create()` stores the ID of the created thread in the location referenced by `thread`. In our implementation, the second argument (`attr`) shall always be `NULL`. The thread is created and executes `start_routine` with `arg` as its sole argument. If the `start_routine` returns, the effect shall be as if there was an implicit call to `pthread_exit()` using the return value of `start_routine` as the exit status. Note that the thread in which `main()` was originally invoked differs from this. When it returns from `main()`, the effect shall be as if there was an implicit call to `exit()` using the return value of `main()` as the exit status.

```
void pthread_exit(void *value_ptr);
```

The `pthread_exit()` function terminates the calling thread. In our current implementation, we ignore the value passed in as the first argument (`value_ptr`) and clean up all information related to the terminating thread. The process shall exit with an exit status of 0 after the last thread has been terminated. The behavior shall be as if the implementation called `exit()` with a zero argument at thread termination time.

```
pthread_t pthread_self(void);
```

The `pthread_self()` function shall return the thread ID of the calling thread. For more details about error handling, please refer to the respective man pages.

Example Library Usage

```
static void* my_thread(void* my_arg) {
    pthread_t my_tid = pthread_self();
    // Do something in this thread
    if (i_need_to_exit_now) {
        pthread_exit(NULL); // Can exit with another value
    }
    // Maybe do some more things in this thread
    // Suggestion: do something that takes a long time
    return NULL; // Can return other values
}

int main(void) {
    pthread_t tid;
    int error_number = pthread_create(
        &tid, NULL, my_thread, NULL /*my_arg*/);
    pthread_t my_tid = pthread_self();
    // my_tid must not equal tid (a different thread)

    // Suggestion: do something that takes a long time, and test that
    // both this and the other long work get to take turns.
}
```

Additionally, see the given `tests/busy_threads.c`. It is an incomplete test case (it may still pass when your code is broken). See its comments for suggested improvements.

Given Code:

The student environment has been updated. Run `docker pull`

`docker.io/dannosliwcd/ec440-student` to access the new tools (clang compiler, clang-tidy, and cppcheck). Since valgrind does not work well with this assignment, those tools provide additional static analysis options you can use to reduce the tooling gap.

There may be a docker bug that claims you need a dockerhub account to download these updates. That is not true. If it tells you to log in you can (if you happen to have an account), or you can delete an existing docker config file (only if it is not applicable for you). In the VM, run `rm ~/.docker/config.json` to delete the existing config file (which in my case was empty, but causing login problems).

We are providing a template that you can optionally use in your solution to this assignment. The template includes empty implementations of the required functions, annotated with comments that outline a possible approach to solving the challenge.

https://drive.google.com/file/d/19Yy6yUixKJjA6NXVlwDi8_NiusUWbCth/view?usp=sharing

Note: if you encounter “longjmp causes uninitialized stack frame” during your intermediate solutions, you may want to add `-U_FORTIFY_SOURCE` to your makefile’s CFLAGS.

Hints:

- 1) You will probably need to use a data structure that can store information about the state of the thread (its set of registers), its stack (e.g., a pointer to the thread's stack area), and status of the thread (whether it is running, ready to run, or has exited). This data structure is often referred to as a *thread control block* (TCB). This is the same concept as Process Control Blocks (PCB), used to manage multiple processes. As you may need to accommodate multiple threads at the same time (you can assume a maximum of 128 threads can be created), the thread control blocks should be stored in a list or a table (array).
- 2) You should create a helper function that initializes your thread subsystem when the application calls `pthread_create()` for the first time. Before the call to the helper function, there is only one thread running (the main program).
- 3) Be sure to `#include <pthread.h>` in your program. That header provides declarations for the functions you are required to define. That is the header our tests will compile against.
- 4) The process of picking another thread is called scheduling. You may cycle through the available threads in a round robin fashion, giving an equal, fair share to each thread.
- 5) In order to switch between threads, the currently executing thread needs to call the `setjmp()` library function. `setjmp` will save the current state of the thread into a `jmp_buf` structure. Then your thread subsystem can pick another thread, use `longjmp()` function along with saved `jmp_buf` to restore to a previously saved state and resume the execution of the new thread.

- a) Think about how often you want the thread of a program to call `setjump` (and thus, give up control of the CPU). Application developers will not do it for you. To solve this issue, you should employ signals and alarms. We can use the `ualarm` or the `setitimer` function to set up a periodic timer that sends a `SIGALRM` signal every X milliseconds (assume that X=50ms for this project). Whenever the alarm goes off, the operating system will invoke the signal handler for `SIGALRM`. So, you can install your own custom signal handler that performs the scheduling (switching between threads) for you. To install this signal handler, you should use `sigaction` with the `SA_NODEFER` flag. Otherwise (e.g., when using the deprecated `signal` function), alarms are automatically blocked while you are running the signal handler.
- b) We **require** that your thread system supports thread preemption and switches between multiple threads that are ready. It is **not okay** to run each individual thread to completion before giving the next one a chance to execute.
- c) To create a new thread, the system has to properly initialize the TCB for the new thread: create a new thread ID, allocate a new stack (`malloc` can come in handy) of 32,767 byte size and initialize the thread's state so that it "resumes" execution from the start function that is given as argument to the `pthread_create` function. For this, we could use `setjump` to save the state of the current thread in a `jmp_buf`, and then, modify this `jmp_buf` in two important ways. First, we want to change the program counter (the RIP) to point to the `start_thunk` function we provide in `ec440threads.h`. Second, we want the stack pointer (the RSP) to point to the top of our newly allocated stack. `start_thunk` is a helper function that moves the value stored in R13 to RDI (why? remember the AMD64 calling convention) and then jumps to whatever address is stored in R12. Hence, for `start_thunk` to work we should store the value of `*arg` in R13 and the address of `start_func` in R12 (note that both R12 and R13 are members of `jmp_buf`).
- d) To modify the `jmp_buf` directly, we have to first understand that it is a very operating system and processor family-specific data structure that is typically not modified directly. `libc` defines the following constants as the eight integer elements of this structure:

```
#define JB_RBX 0
#define JB_RBP 1
#define JB_R12 2
#define JB_R13 3
#define JB_R14 4
#define JB_R15 5
#define JB_RSP 6
#define JB_PC 7
```

- i) How can you know this? It depends on internal architecture-specific data structures in `libc`. We can call `gnu_get_libc_version()` to see that we have `glibc 2.27` in the student environment. We don't have the `glibc` internal headers pre-installed in our student environment, but here's a view for our

version, in x86_64 architectures.

https://elixir.bootlin.com/glibc/glibc-2.27/source/sysdeps/x86_64/jmpbuf-of-fsets.h

- e) We can see that the stack pointer has index 6 and the program counter has index 7 into the `jmp_buf`. This allows us to easily write the new values for RSP and RIP into a `jmp_buf`. Unfortunately, there is a small complication on the Linux systems in the student/grading environment. These machines are equipped with a libc that includes a security feature to protect the addresses stored in jump buffers. This feature "mangles" a pointer before saving it in a `jmp_buf`. To convince yourself of this, generate a `jmp_buf` while stepping through the program with a debugger, and print the contents in the `jmp_buf`. You will see that many register values are exactly as in the debugger, but not RIP and RSP. Thus, we also have to mangle our new stack pointer and program counter before we can write it into a jump buffer, otherwise decryption (and subsequent uses) will fail. `longjmp` will automatically decrypt these values, so it is imperative that they are properly encrypted/mangled beforehand. To mangle a pointer before writing it into the jump buffer, make use of the following function:

```
unsigned long int ptr_mangle(unsigned long int p){
    unsigned long int ret;
    asm("movq %1, %%rax;\n"
        "xorq %%fs:0x30, %%rax;"
        "rolq $0x11, %%rax;"
        "movq %%rax, %0;"
        : "=r"(ret)
        : "r"(p)
        : "%rax"
    );
    return ret;
}
```

- f) Remember that the start routine of every new thread expects a single argument (a void pointer called `arg`). When a new thread is launched, the start routine should run as if it were invoked by a preceding call instruction (but it isn't, we're setting up a "fake" context with the `jmp_buf` and the stack). Hence, we have to basically "simulate" such a function call to the start routine. To this end, we just have to initialize the new stack for a thread and pass the argument according to the calling convention. Our environment uses the amd64 calling convention, hence the first six arguments are passed in registers. To make this step easier, we provide the `start_thunk` function that behaves as described above. Specifically, whatever value is stored in R13 (`*arg`) when `start_thunk` is called, will be passed on as the first argument to the function referenced in R12 (`start_routine`). For completeness sake, `start_thunk` is defined as:

```
void *start_thunk() {
    asm("popq %%rbp;\n" //clean up the function prologue
        "movq %%r13, %%rdi;\n" //put arg in $rdi
        "pushq %%r12;\n" //push &start_routine
    );
    return start_routine;
}
```

```

        "retq;\n"                //return to &start_routine
        :
        :
        : "%rdi"
    );
    __builtin_unreachable();
}

```

In addition to the argument, we also need to make sure that the new stack has a correct return address, such as in struct `start_routine`, to return to after completion. Given that the specification requires that `start_routine` invokes `pthread_exit` upon completion, we can easily achieve this by putting the address of `pthread_exit` on the top of the stack, which will transparently be treated by `start_routine` as the return address. Make sure to adjust the stack pointer (RSP) appropriately.

- g) Once your stack is initialized and the (mangled) stack pointer (RSP) and program counter (RIP) are written to the `jmp_buf`, your new thread is all ready to go, and when the next SIGALRM arrives, it is ready to be scheduled!
- h) In your solution you will have to use `ptr_mangle` to store the mangled stack and instruction pointers into the `jmp_buf` structure. If, during development, you want to inspect these values (e.g., to check you actually got a true `jmp_buf`) you can use the following counterpart to the above `ptr_mangle` function (not necessary, just for convenience)

```

unsigned long int ptr_demangle(unsigned long int p) {
    unsigned long int ret;
    asm("movq %1, %%rax;\n"
        "rorq $0x11, %%rax;"
        "xorq %%fs:0x30, %%rax;"
        "movq %%rax, %0;"
        : "=r"(ret)
        : "r"(p)
        : "%rax"
    );
    return ret;
}

```

- 6) Unfortunately, you will **not be able to use valgrind or sanitizers** with this assignment. Those tools are not aware of our `jmp_buf` modifications to replace the active stack, so they are likely to raise false alarms. Static analysis tools will still be helpful, though. As always, do not ignore compiler warnings! Try running `make static_analysis!`

Submission Guidelines:

Your threading library must be written in C and run on the Linux Gradescope environment.

Submit your **source code and makefile** to Gradescope. When we run `make` in your submission,

a threads.o object file must be produced, containing the compiled definitions of the required functions listed in this homework prompt. This file must not define a main function.

Your threading library **must not** depend on libpthread. You are implementing an alternative to libpthread in this assignment.

Your final submission must also include a `README` file (README.md is also okay). If you relied on any **external resources** to help complete this assignment, note the resource and the challenge it helped resolve. **If not**, say so. Use this file to provide any **additional notes** you would like to share with instructors about your submission. Aside from that, it just needs a **short description** of the purpose of the project. But you are permitted to include additional details that *you* want in your README.

Oral Exams

When the submission deadline is reached, we will start oral exam sessions over Zoom. In these sessions, we will ask you to explain how some parts of your program work. Details about how to schedule an oral exam time will be posted in the week leading up to the exams.