

# **EC 440 – Introduction to Operating Systems**

**Orran Krieger (BU)**  
**Larry Woodman (Red Hat)**

# Computer Security

**Security is a big field, encompassing:**

- Application security
- Network security
- Authentication
- Digital forensics
- Cryptography
- Privacy/Anonymity
- etc.

# Operating Systems Security

- In this course we are only going to concern ourselves with security as it applies to operating systems
- This is some mix of *authentication*, *access control*, and *application security*
- First though, some definitions and ideas about how to think about security

# Computer Security

**Generally, we talk about computer security in three broad categories:**

- *Confidentiality* – preventing others from finding out information we don't want them to have
- *Integrity* – preventing others from modifying our data without permission
- *Availability* – preventing others from denying us access to some service (denial of service)

# Threat Modeling

- It usually doesn't make sense to talk about a system being "secure" or "unsecure"
- Instead, we need to be more precise:
  - What are we trying to protect?
  - Who do we need to protect against? What are their capabilities?

# A Practical Threat Model

## **Threat:**

- Ex-girlfriend/boyfriend breaking into your email account and publicly releasing your correspondence with the My Little Pony fan club

## **Solution:**

- Strong passwords

# A Practical Threat Model

## **Threat:**

- Organized criminals breaking into your email account and sending spam using your identity

## **Solution:**

- Strong passwords + common sense (e.g., don't click on unsolicited herbal Viagra ads that result in keyloggers and sorrow)

# A Practical Threat Model

## **Threat:**

- The Mossad doing Mossad things with your email account

## **Solution:**

- Magical amulets(cross, voodoo doll)?
- Fake your own death, move into a submarine?
- YOU'RE STILL GONNA BE MOSSAD'ED UPON



# Threat Modeling

**Roughly, we can divide this into three steps:**

- System understanding
- Threat categorization
- Countermeasures and mitigation


# System Understanding

**To properly protect a system, you need to understand it:**

- Identify assets that need protecting
- Look at ways the system can be used and assets can be accessed
- Figure out what right should be granted to what assets and classes of users
- Identify *privilege boundaries* – places where a program or user changes their privilege level

# Example System Diagram

## Legend



External  
Entity



Data Store



Process



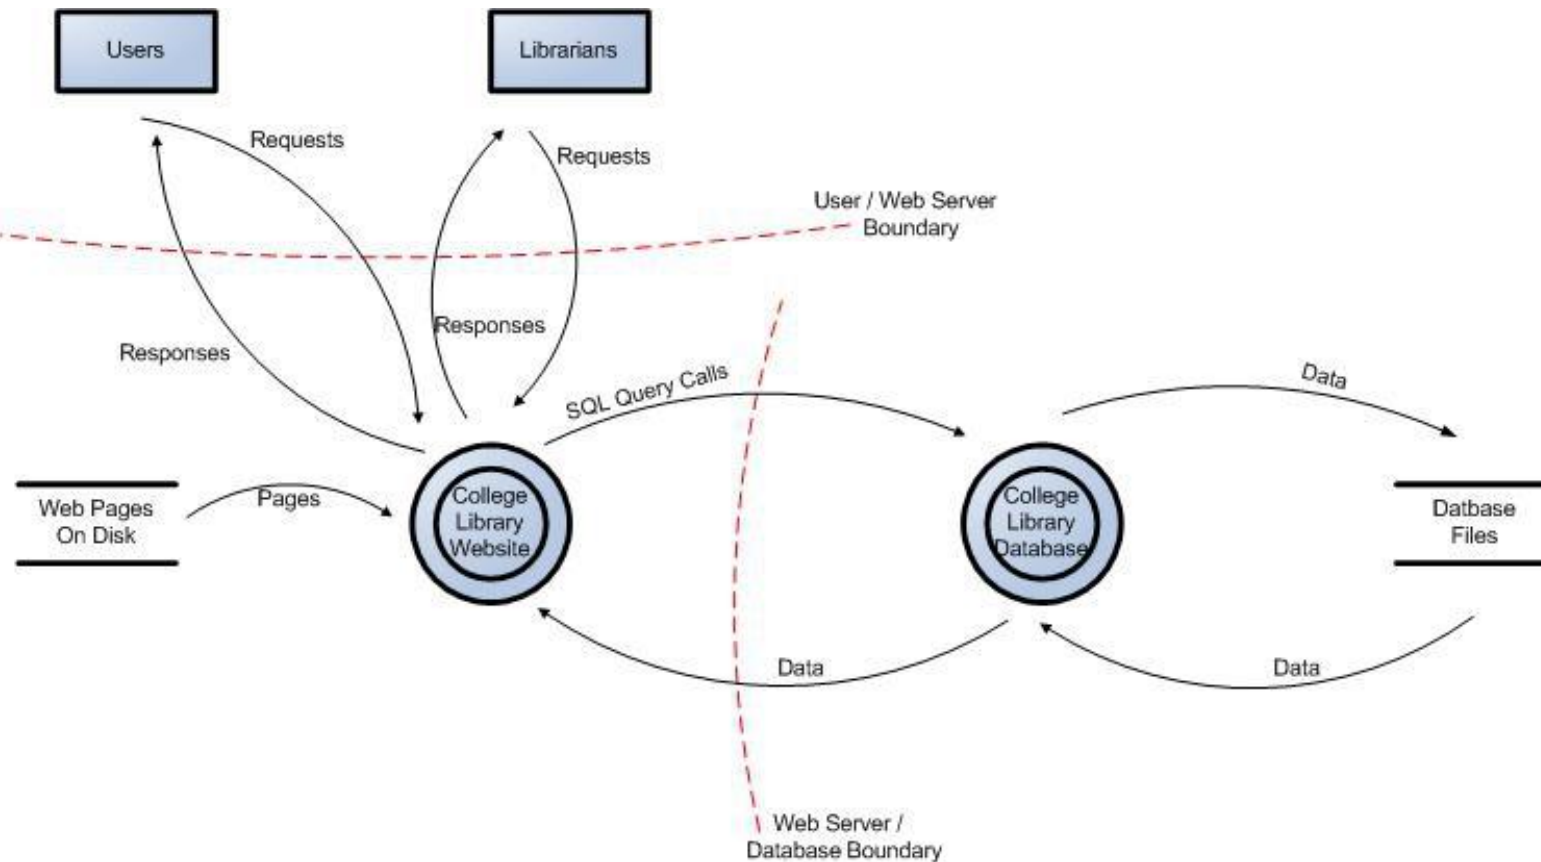
Multiple  
Process



Data Flow



Privilege Boundary



Source:

# Threat Categorization

- Look at the system from an attacker's point of view
- What goals might an attacker have?
- How could they achieve these goals?
- May help to use a threat categorization such as **STRIDE**: **S**poofing, **T**ampering, **R**epudiation, **I**nformation Disclosure, **D**enial of Service, **E**levation of Privilege

# Countermeasures and Mitigation

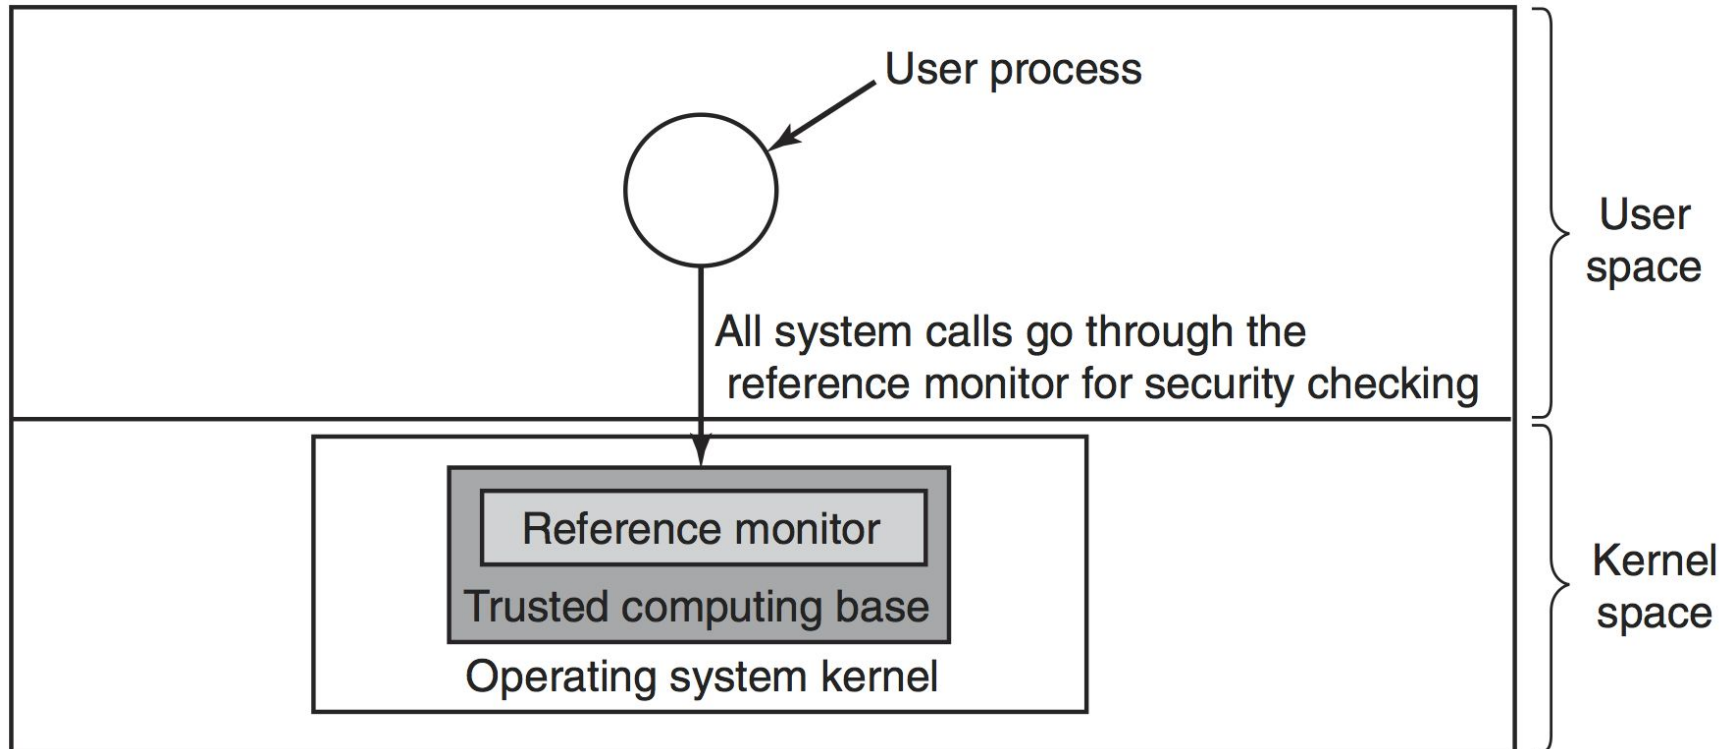
**For each of the threats to some asset, come up with a plan for mitigating or nullifying the threat, for example:**

- Attacker might guess someone's password
  - -> enforce password complexity requirements
- Attacker might snoop on network traffic
  - -> encrypt data that is sent on the network

# Trusted Computing Base(TCB)

- One strategy for building secure operating systems is to organize them into *trusted* and *untrusted* components
- The goal is that if the *trusted* components performs according to its specification, then some specific set of guarantees about security must hold
- A *reference monitor* checks all accesses between trusted and untrusted components

# TCB + Reference Monitor



# Aside: Bugs & Program Size

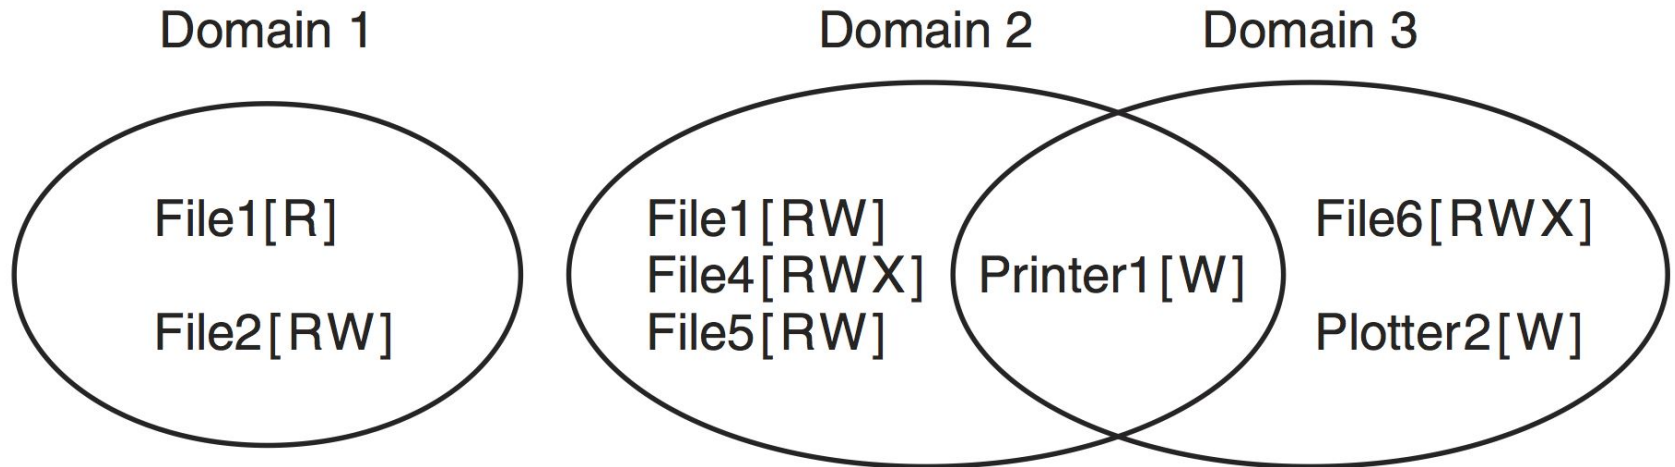
- One rule of thumb is that as program size increases, the number of bugs increases as well
- It's harder to reason about code the more of it there is and the more complex it is.
- Therefore, if we want to be confident in our trusted computing base, we should work to *minimize* the amount of code in it
- How big is the TCB for widely-used operating systems?(HUGE!)



# Protection Domain

- A protection domain is a set of (object, rights) pairs
- A right, in this context, means an operation that can be performed on an object
- So, for example, a protection domain might correspond to a user and the set of objects they have access to
- Or, a group of users that all share the same rights

# Protection Domains



# Principle of Least Privilege

- One principle for designing secure systems is the *principle of least privilege*.
- The core idea is that the set of objects and rights for a given protection domain should be as small as possible
- This seems obvious, but is often violated in practice

# Processes and Protection Domains

- At any given time, a process operates in one protection domain
  - i.e., there is some specific set of objects that the process has permissions to perform some actions on
- Processes can also typically switch between protection domains as they run
  - The rules for when and how they do this vary widely between different operating systems

# The UNIX Protection Model

- The protection domain of a process in UNIX is defined by its user id (uid) and group id (gid)
- The objects are files (including special files like hardware devices)

# The UNIX Protection Model

- Each process is further divided into two halves:
  - kernel mode and user mode
- The kernel half can access a different set of objects from the user half, so the change from user to kernel is a domain switch
- Executable files can also have SETUID and SETGID bits, meaning that when they are executed they will run under different permissions

# User Management

**Code running in user mode is always linked to a certain identity**

- Security checks and access control decisions are based on user identity

**Unix is user-centric**

- No roles
- Users Identified by user id (uid) & group id (gid)
- Authenticated by password (stored encrypted)

**User root**

- Superuser, system administrator
- Special privileges (access resources, configure the OS)
- Cannot decrypt user passwords

# Process Management (PCB)

## Process Attributes

- Process id (PID)
  - Uniquely identifies a process
  - PIDs are reused
- User id (UID)
  - ID of owner of process
- Effective user id (EUID)
  - ID used for permission checks / access control
- Saved user id (SUID)
  - To temporarily drop and restore privileges
- Lots of management information
  - Scheduling
  - Memory management
  - Resource management



# User Authentication

## How does a process get a user ID?

- Authentication via login

## Passwords

- User passwords used as key for crypt() function
- 12 bit public salt (prevent pw collisions)
- Repeatedly apply encryption

## Password cracking

- Dictionary attacks
- Crack, JohnTheRipper

# User Authentication

## **File /etc/passwd**

- Maps user names to user ids (many applications legitimately need this)
- No legitimate need for encrypted passwords

## **File /etc/shadow**

- Contains encrypted passwords
- Account information (last change, expiration)
- Readable only by superuser and privileged programs
- Different hash algorithms supported
  - DES
  - MD5
  - SHA-{256,512}

# Unix Groups

## **Users belong to one or more groups**

- Primary group (stored in /etc/passwd)
- Additional groups (stored in /etc/group)
- Possibility to set group password
- Become group member with newgrp

## **File /etc/group**

groupname : password : group id : additional users

root:x:0:root

bin:x:1:root,bin,daemon

users:x:1000:pizzaman

## **Special group wheel**

- Group for users that can call su

# File System

## Access Control

- Permission bits
- chmod, chown, chgrp, umask
- File listing

-                      rwx              rwx                      rwx  
(file type) (user) (group) (other/world)

Type	r	w	x	s	t
File	Read access	Write access	Execute	suid / sgid inherit id	sticky bit
Directory	List files	Add and remove files	Stat / execute files, chdir	New files have dir-gid	Files only deletable by owner

# SUID Programs

**Each process has *real* and *effective* user / group id**

- Usually identical
- Real id
  - Determined by current user
  - login, su
- Effective ids
  - Determine access rights of a process
  - System calls (e.g., `setuid()`)
- `suid/sgid` bits
  - Start process with effective ID different from real ID
  - Attractive targets for attacker

**No SUID shell scripts anymore**

# Extended Attributes

```
# lsattr /etc/passwd /etc/ssl
```

```
-----e-- /etc/passwd
```

```
-----I--e-- /etc/ssl/certs
```

- Require support from file system
- Management via `lsattr`, `chattr`
  - Undeleteable (u)
  - Append only (a)
  - Immutability (i)
  - Secure deletion (s)
  - Compression (c)
  - Hashed trees indexing for directories (I)

# POSIX ACLs

**Extend UNIX permission model to support fine-grained access control**

```
$ sudo setfacl -m u:pizzaman:r secret
$ getfacl secret
# file: secret
# owner: root
# group: root
user::rw-
user:pizzaman:r--
group:---
mask:r--
other:---
```

# Software Security

- One final aspect to OS security is *software security* – identifying and preventing programming flaws that could let an attacker take control
- Many of these are caused by the fact that languages currently in use are not *memory safe* – it is possible to write to data outside of program variables
- A big offender here is C/C++



# Operating System Defenses

- OSes can be designed to make problems in user-level applications harder to exploit
- These generally don't make attacks on software impossible, but they can make them much more difficult
- This is something of an arms race – defenders come up with new mechanisms, attackers find ways around them

# Classic Problem

```
#include<stdio.h>
#include<string.h>

int main(int argc, char ** argv) {
    char buf[256];
    strcpy(buf, argv[1]);
    printf("%s\n", buf);
    return 0;
}
```

What can possibly go wrong?

# Stack Buffer Overflow

- Recall the standard stack frame:
  - Local variables
  - Return address
- If we try to store too much data in a local stack variable (e.g., a character array) we will overwrite the return address
- When the ret instruction is executed, control will jump to somewhere controlled by user input

# Stack Canaries / Cookies

- Idea: put a special value in between the local variables and the return address so that overflowing a local buffer can be detected
- Upon entering the function, set a randomly-generated *cookie* value on the stack and store a backup copy somewhere
- Before executing a ret, check the stack cookie value against the backup and raise an error if it fails

# PCB In Linux – task\_struct

```
struct task_struct {  
    volatile long state; /* -1 unrunnable, 0 runnable, >0 stopped */  
    /* task state */  
    int exit_state;  
    pid_t pid;  
    /* Canary value for the -fstack-protector gcc feature */  
    unsigned long stack_canary;  
    struct task_struct *parent; /* recipient of SIGCHLD, wait4() reports */  
    struct timespec start_time; /* monotonic time */  
    char comm[TASK_COMM_LEN]; /* executable name excluding path */  
    /* CPU-specific state of this task */  
    struct thread_struct thread;  
    /* signal handlers */  
    struct signal_struct *signal;  
    sigset_t blocked, real_blocked;  
    ... 300 lines ...  
}
```

# What's the Problem?

Attack can *inject* new code and then corrupt the program to *execute* this code ....

What's the solution?

– Make it so that:

1. The attacker cannot inject new code, or  
(Hard to do for programs written in C/C++)
2. The new code cannot be executed!

# DEP/NX/W $\oplus$ X

- Another defense is to try and make sure that even if an attacker can overflow a buffer and change the return address, they cannot execute the injected code
- In the previous example, attacker code was placed into a stack buffer
- So, simple solution: don't allow data to be executable!
- Generally this requires hardware support
  - NX bit in x86 page protections





# Current State of Affairs

- Great, we just made sure that the attacker can no longer execute any code that he injected into the program!
- Are we done?
- If the attacker cannot run injected code what else can he do?

Reuse existing code (which must be executable as it is code) → Code Reuse Attacks

# Code Reuse Attacks

- Despite DEP, attacks can still run code!
- Instead of trying to execute their own code, attackers can change the return address to point to existing code in memory
- By setting up values on the stack, they can bounce around executing tiny snippets of code ending in ret
- Thus, by chaining these together, arbitrary computation can be performed – without executing anything marked as data

# Code Reuse Attacks

- Problem: Attacker *knows where* existing code is in memory and manages to run that code
- What can we do to thwart such attacks?
- Make it so the attacker does no longer know where code is in memory
- How would we do that?
- We make sure the code is loaded at random addresses every time a new process starts (fork, exec, both?)

# Address Space Layout Randomization

- Exploiting a buffer overflow typically requires knowing about the precise layout of memory
- For example, we may need to know where the stack is located, or where a certain library has been loaded
- Thus, to make attackers' lives more difficult, we can place the program, stack, and libraries at random locations each time the program starts

# ASLR and Address Space

- We can estimate the amount of randomness provided by ASLR by counting the number of possible locations to load things
- On a 32-bit system, address space is not very large, so there are not many ways to randomize
- In 2004, researchers showed that on 32-bit systems, ASLR only has about 16 bits of entropy (65,536 possible values)
- The correct location can be guessed in a matter of seconds by just trying each possibility

# Side Channel Attacks

- Sophisticated techniques that measure performance/timing differences to determine the contents of caches and therefore memory.
- Fix required hardware updates and/or firmware/kernel modifications that usually result in performance degradation.
- Meltdown - sharing page table between user and kernel results in security holes
- Spectre - hardware prefetch and speculation results in security holes