

# **EC 440 – Introduction to Operating Systems**

**Orran Krieger (BU)  
Larry Woodman (Red Hat)**

# Outline

- What is an operating system
- How we got here
- Our background
- What you will learn & course administration

# What is an Operating System?

Application

Computer

# Problem: HW interfaces are ugly

Application

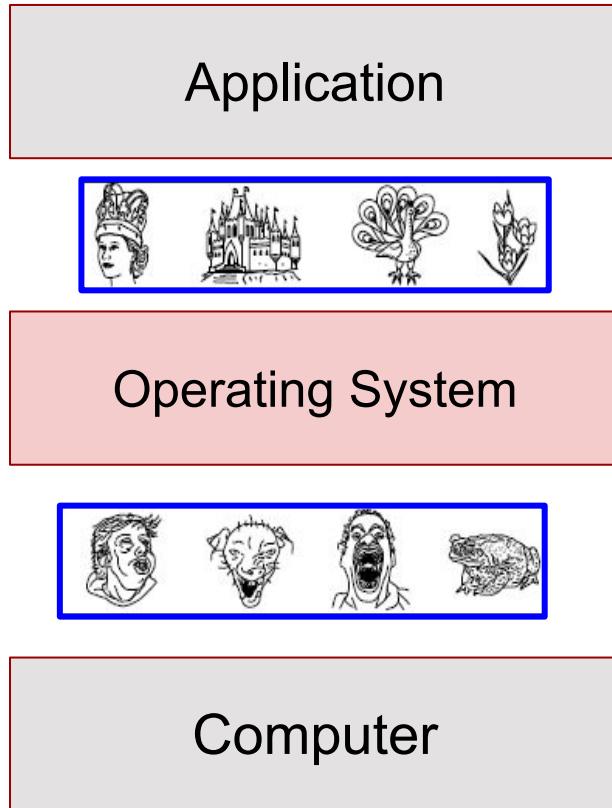


Computer

- CPUs
- Memory
- Caches
- Buses
- Storage devices
- Serial ports
- Network interface cards

If you needed to understand all the Hardware, no code would ever be written.

# At base level, the OS creates clean interfaces



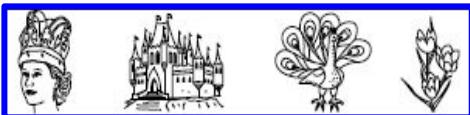
- Processes/threads
- Directories/Files
- Virtual memory
- Sockets
- ...

# Allow many applications to run

Application

Application

Application



Operating System



Computer

- Ensuring they can't break each other.
- Meeting their needs:
  - real time, e.g., car, gaming...
  - CPU intensive, e.g. game, computational science

# From Different Users



Application

Application

Application



Operating System

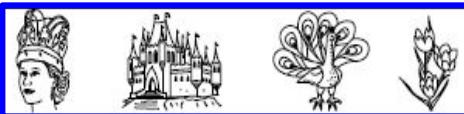


Computer

- Protecting users from each other
- Ensuring that each gets their fair use of resources

# On many different types of computers

Application



Operating System

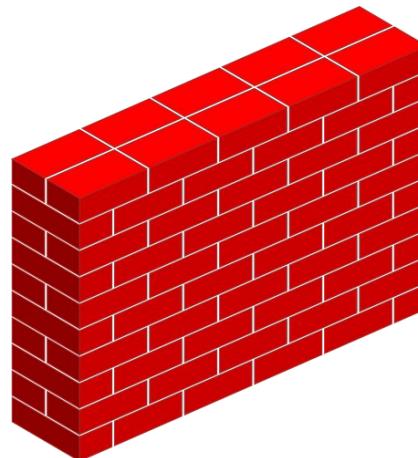
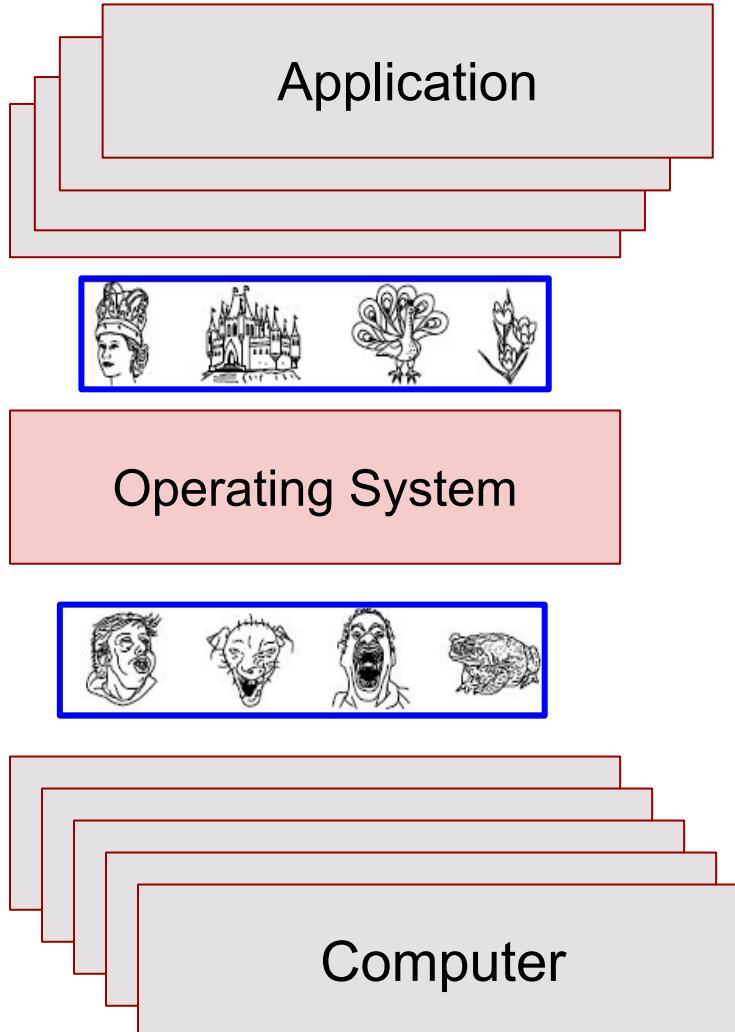


Computer

- Processor Architectures: x86, Power, ARM, RISC5...
- Across: PCs, servers, cars, phones, tablets, ...

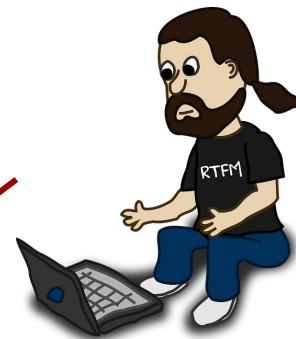
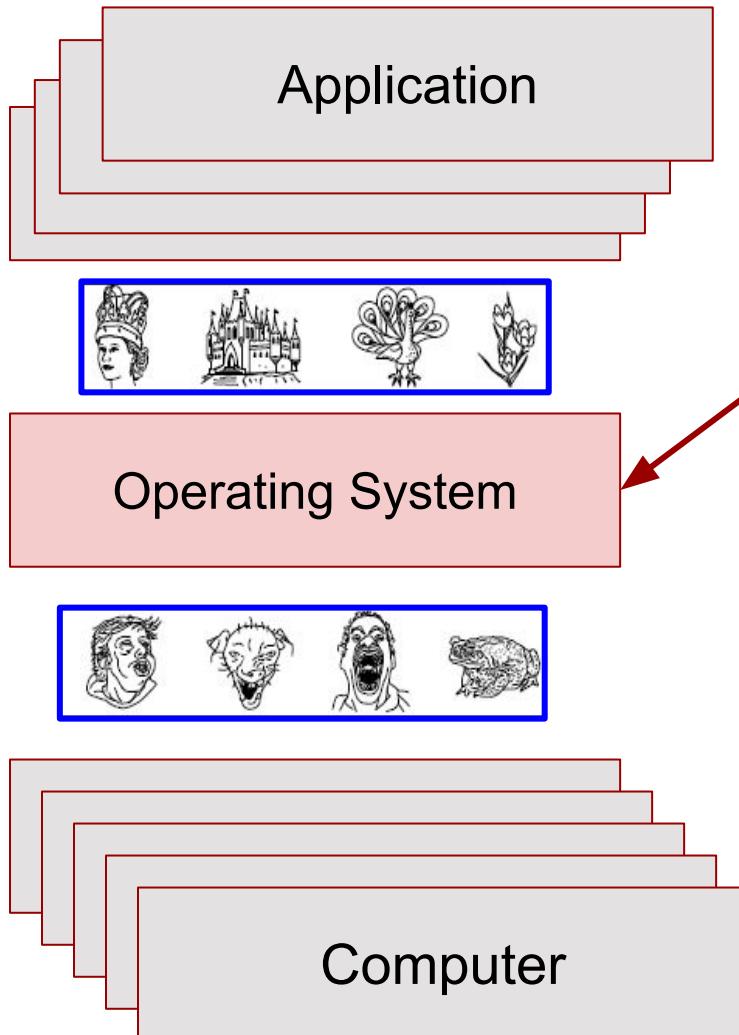
Enabling common applications and tools...

# Protecting against attacks



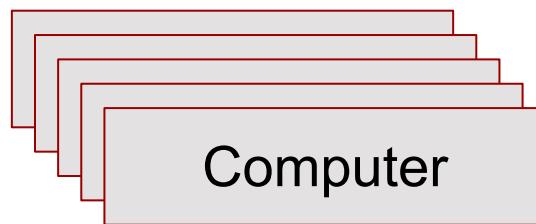
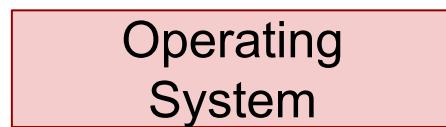
- Firewalls
- Logging
- Security policies

# Enabling configuration control, automation, upgrading...

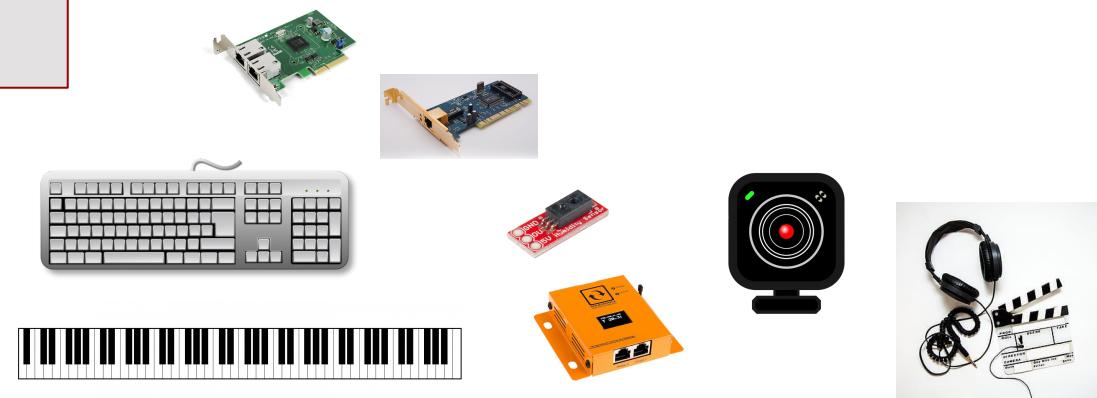
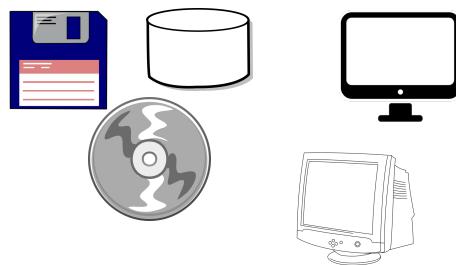


- Personal Computers
- Embedded devices: cars, fridges, routers...
- Massive cloud data centers
- High Performance Computing environments

# Enabling devices



- Huge range of devices: disks, monitors, keyboards NICs, sensors...
- Manufacturer needs to work with OS
- OS enabling them to be used across many different computers

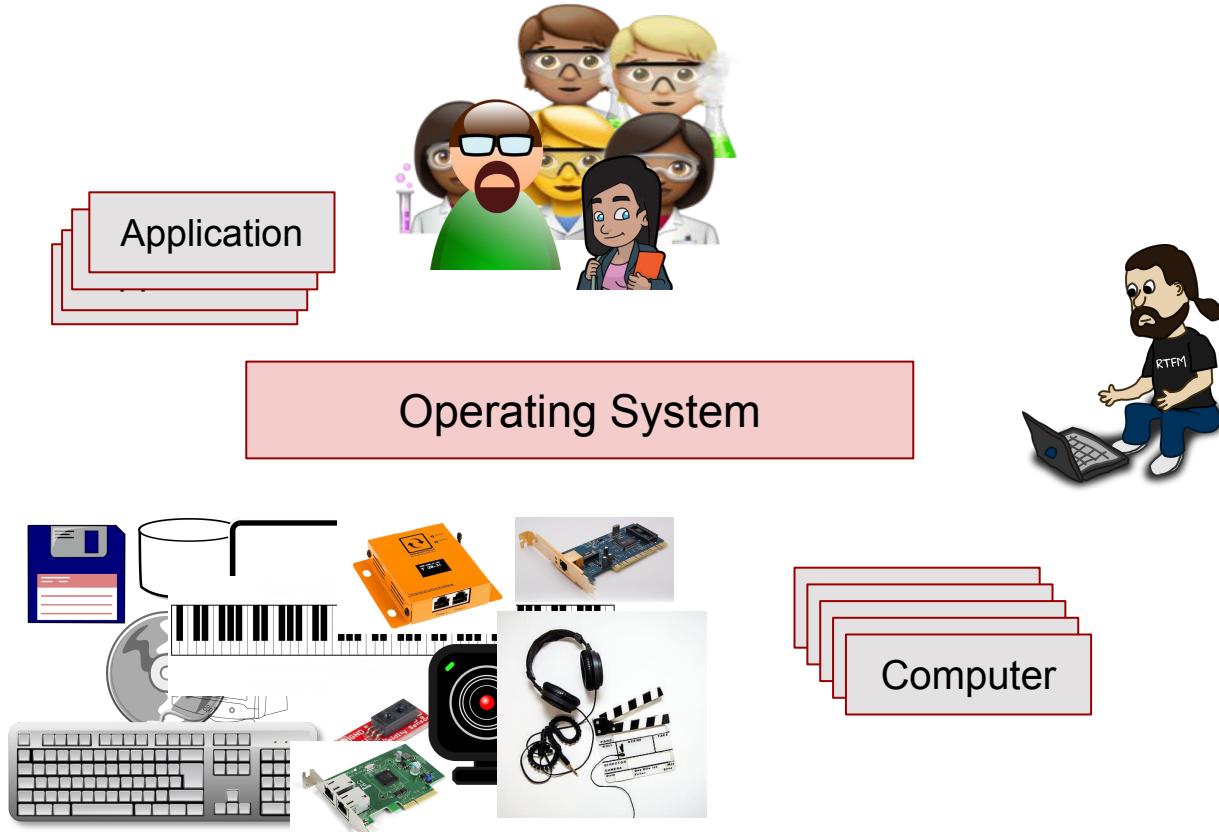


# So the job of the modern OS is to

- Create abstractions on computer resources that abstract those resources
- Manage and multiplex the hardware between different application
- Enabling:
  - application to be run on wide set of computers
  - broad set of devices
  - wide range of use cases: phone to cloud
- All while enforcing security

# Why do we think is so cool?

Core platform, everything relies in it, if you control the platform you control everything.



# This means...

- Focus of some of the largest companies in the world: Microsoft, VMware, Red Hat
- OS people play a key role in every major technology company: Google, Facebook, ..
- Every major change in technology requires OS innovation.
- Led to broader systems, distributed systems: OSDI/SOSP/Usenix are where Borg/Ceph... published
- Started the world of Open Source
- Key to the cloud

# Open Source



- Operating systems started the whole move to Open Source...
- We used to think that only cathedral builders could create a stable OS, but Linux showed BAZAAR could work. it is becoming the dominant OS in cloud, mobile devices, enterprise...
  - Keeps any company from owning monopoly
  - Everyone can scratch their own itch
- Open Source has become requirement for platforms, e.g., HDFS, Hadoop, SPARK, Kubernetes, Ceph...

# Outline

- What is an operating system
- How we got here
- Our background
- What you will learn & course administration

# History of OS/HW from book

Gen 1: 1945–55:

vacuum tubes

Gen 2: 1955–65:

transistors

Gen 3: 1965–1980:

Integrated cir

Gen 4: 1980– : PCs

Gen 5: 1990– :

Mobile

Gen 6: 2015- : Cloud

# Generation 1: Vacuum tubes

- Gen 1: 1940–55:  
vacuum tubes
- Gen 2: 1955–65:  
transistors
- Gen 3: 1965–1980:  
Integrated cir
- Gen 4: 1980– : PCs
- Gen 5: 1990– :  
Mobile
- Gen 6: 2015- : Cloud

- No operating system: all programming done in machine language, or plugboard
- One program & one user at a time on the computer
- Computers special purpose, size of rooms
- Watson, 1945, worldwide market for maybe 5 computers

# Generation 2: Transistors & Batch

- Gen 1: 1945–55:  
vacuum tubes
- Gen 2: 1955–65:  
transistors
- Gen 3: 1965–1980:  
Integrated cir
- Gen 4: 1980– : PCs
- Gen 5: 1990– :  
Mobile
- Gen 6: 2015- : Cloud

- Rise of the mainframe that can be sold to companies
- Only interfaces are pull card deck (fortran) into machine and dump results
- Can quickly run programs one after the other
- First OSes:
  - Job control language
  - Compiler for Fortran

# Generation 3: Transistors & Batch

- Gen 1: 1945–55:  
vacuum tubes
- Gen 2: 1955–65:  
transistors
- Gen 3:** 1965–1980:  
Integrated cir
- Gen 4: 1980– : PCs
- Gen 5: 1990– :  
Mobile
- Gen 6: 2015- : Cloud

- First OS that runs across multiple lines (IBM 360)
- Spooling
- Multiprogramming
- Timesharing
- Virtualization
- MULTICS: Utility computing, hierarchical file system, segments...
- Unix: single-user version of MULTICS
  - MacOS, Linux (Android)
  - Open Source Key success

# Generation 4: Personal Computers

- Gen 1: 1945–55:  
vacuum tubes
- Gen 2: 1955–65:  
transistors
- Gen 3: 1965–1980:  
Integrated cir
- Gen 4:** 1980– : PCs
- Gen 5: 1990– :  
Mobile
- Gen 6: 2015- : Cloud

- Wide variety of OSes for hobbyist... then IBM decides to create PC
- MS-DOS creates massive monopoly, first to recognize OS as control point
- Eventually MAC adopts BSD (& microkernel), BSD license

Networks became ubiquitous, then internet...

# Generation 5: Mobile

- Gen 1: 1945–55:  
vacuum tubes
- Gen 2: 1955–65:  
transistors
- Gen 3: 1965–1980:  
Integrated cir
- Gen 4: 1980– : PCs
- Gen 5:** 1990– :  
Mobile
- Gen 6: 2015- : Cloud

- A new fight for the winning OS... started in 1997
- Now two main competitors, Google's Android which is Linux & iOS
- PCs increasingly irrelevant; world moving to mobile devices and cloud

# Generation 6: Cloud

- Gen 1: 1945–55:  
vacuum tubes
- Gen 2: 1955–65:  
transistors
- Gen 3: 1965–1980:  
Integrated cir
- Gen 4: 1980– : PCs
- Gen 5: 1990– :  
Mobile
- Gen 6: 2015- : Cloud

- In late 90s, Virtualization made a comeback with VMware (more later)
- Suddenly computing became a fungible commodity.
- AWS creates a cloud using Xen; open source virtualization based on Linux

# Is the fun over?

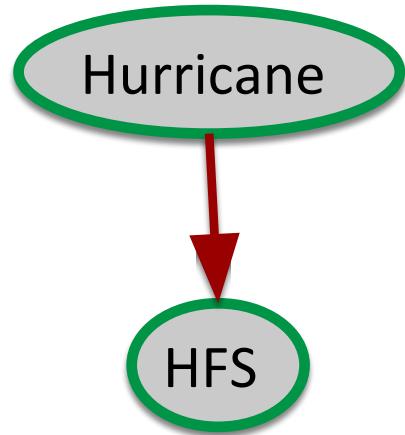
There are fundamental changes in usage and technology:

- Cloud pushing new models of OS for SaaS/FaaS
- Fundamental change in security requirements; provider/tenant
- Denard scaling is over; but network and storage keep getting faster...
- Edge & Cloudlets driven by 5G
- First generation of persistent memory last year
- Rise of smart NICs
- 100 GB networks spanning the world..
- Rise of AI...
- ....

All these require fundamental OS innovation, e.g.: unikernel, partitioning HW, AI controlled policy...

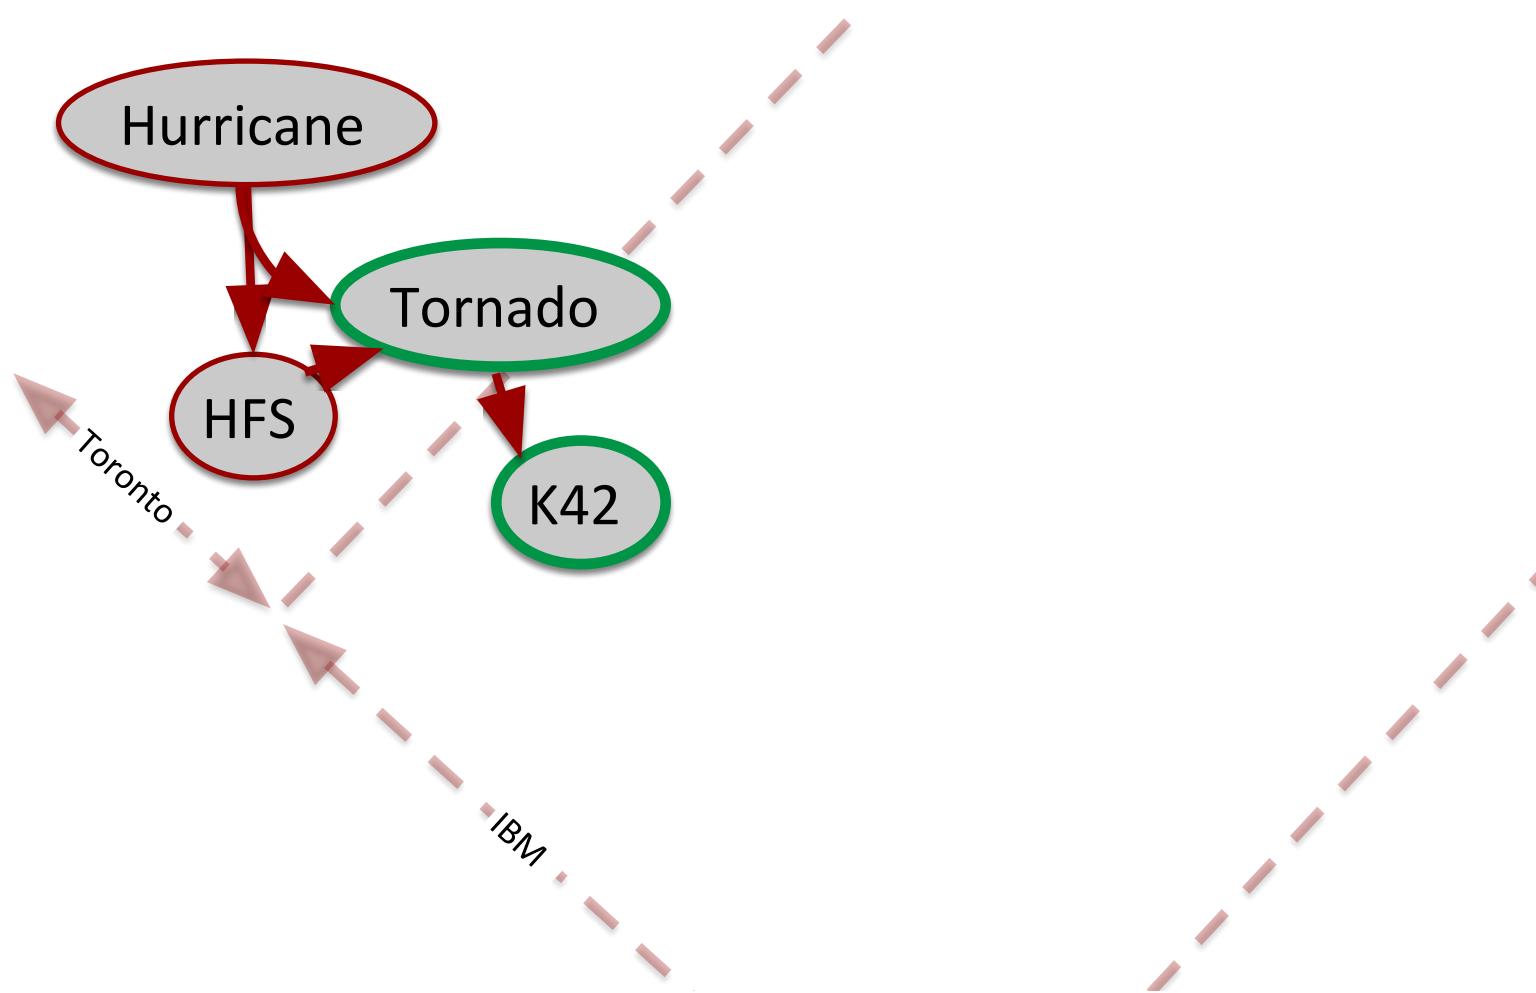
# Outline

- What is an operating system
- How we got here
- Our background
- What you will learn & course administration



## Orran's background

Large-scale NUMA shared memory multiprocessors will become important and current ad hoc approach scalability won't work



Take 2: Large-scale **64-bit** NUMA shared memory multiprocessors will become important and current ad hoc approach scalability won't work. Customizability key to performance

Hurricane

Tornado

HFS

K42

Libra

PS3

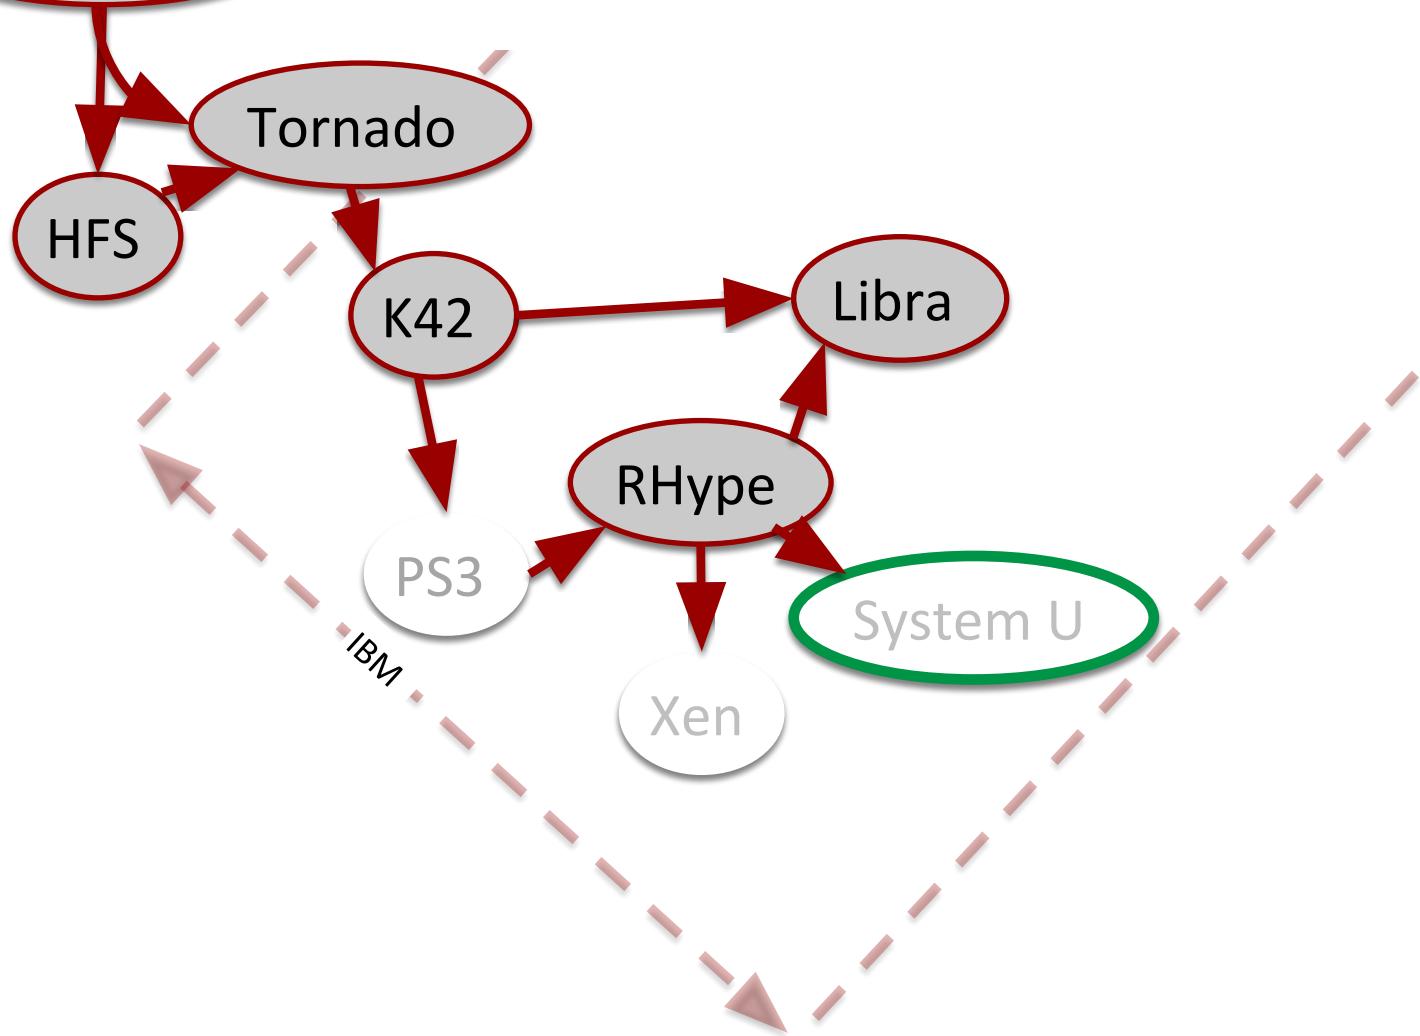
RHype

Xen

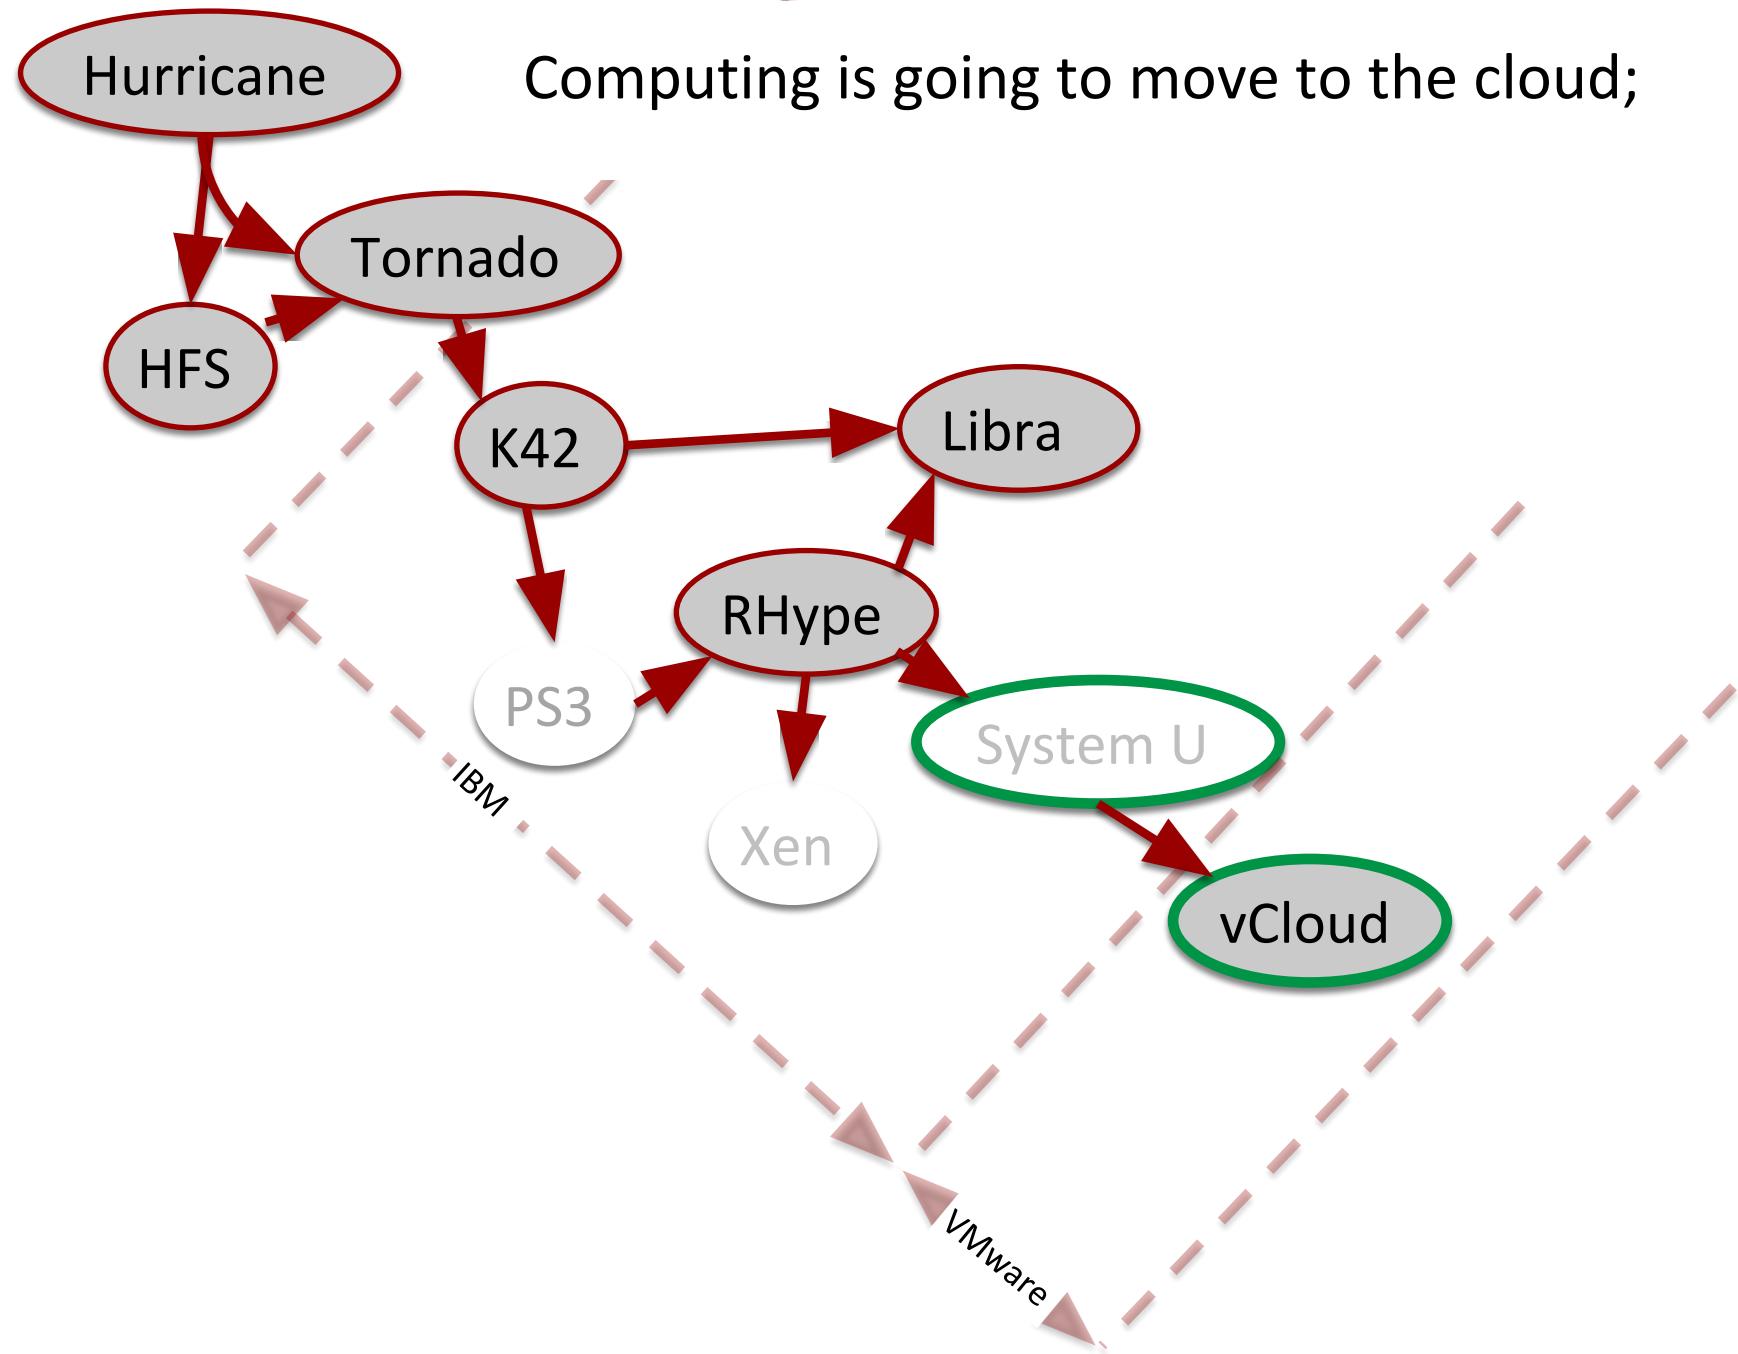
IBM

Virtualization adds fundamental level of indirection:  
deterministic performance, HPC, library Oses for  
applications, for devices...

Computing is going to move to the cloud

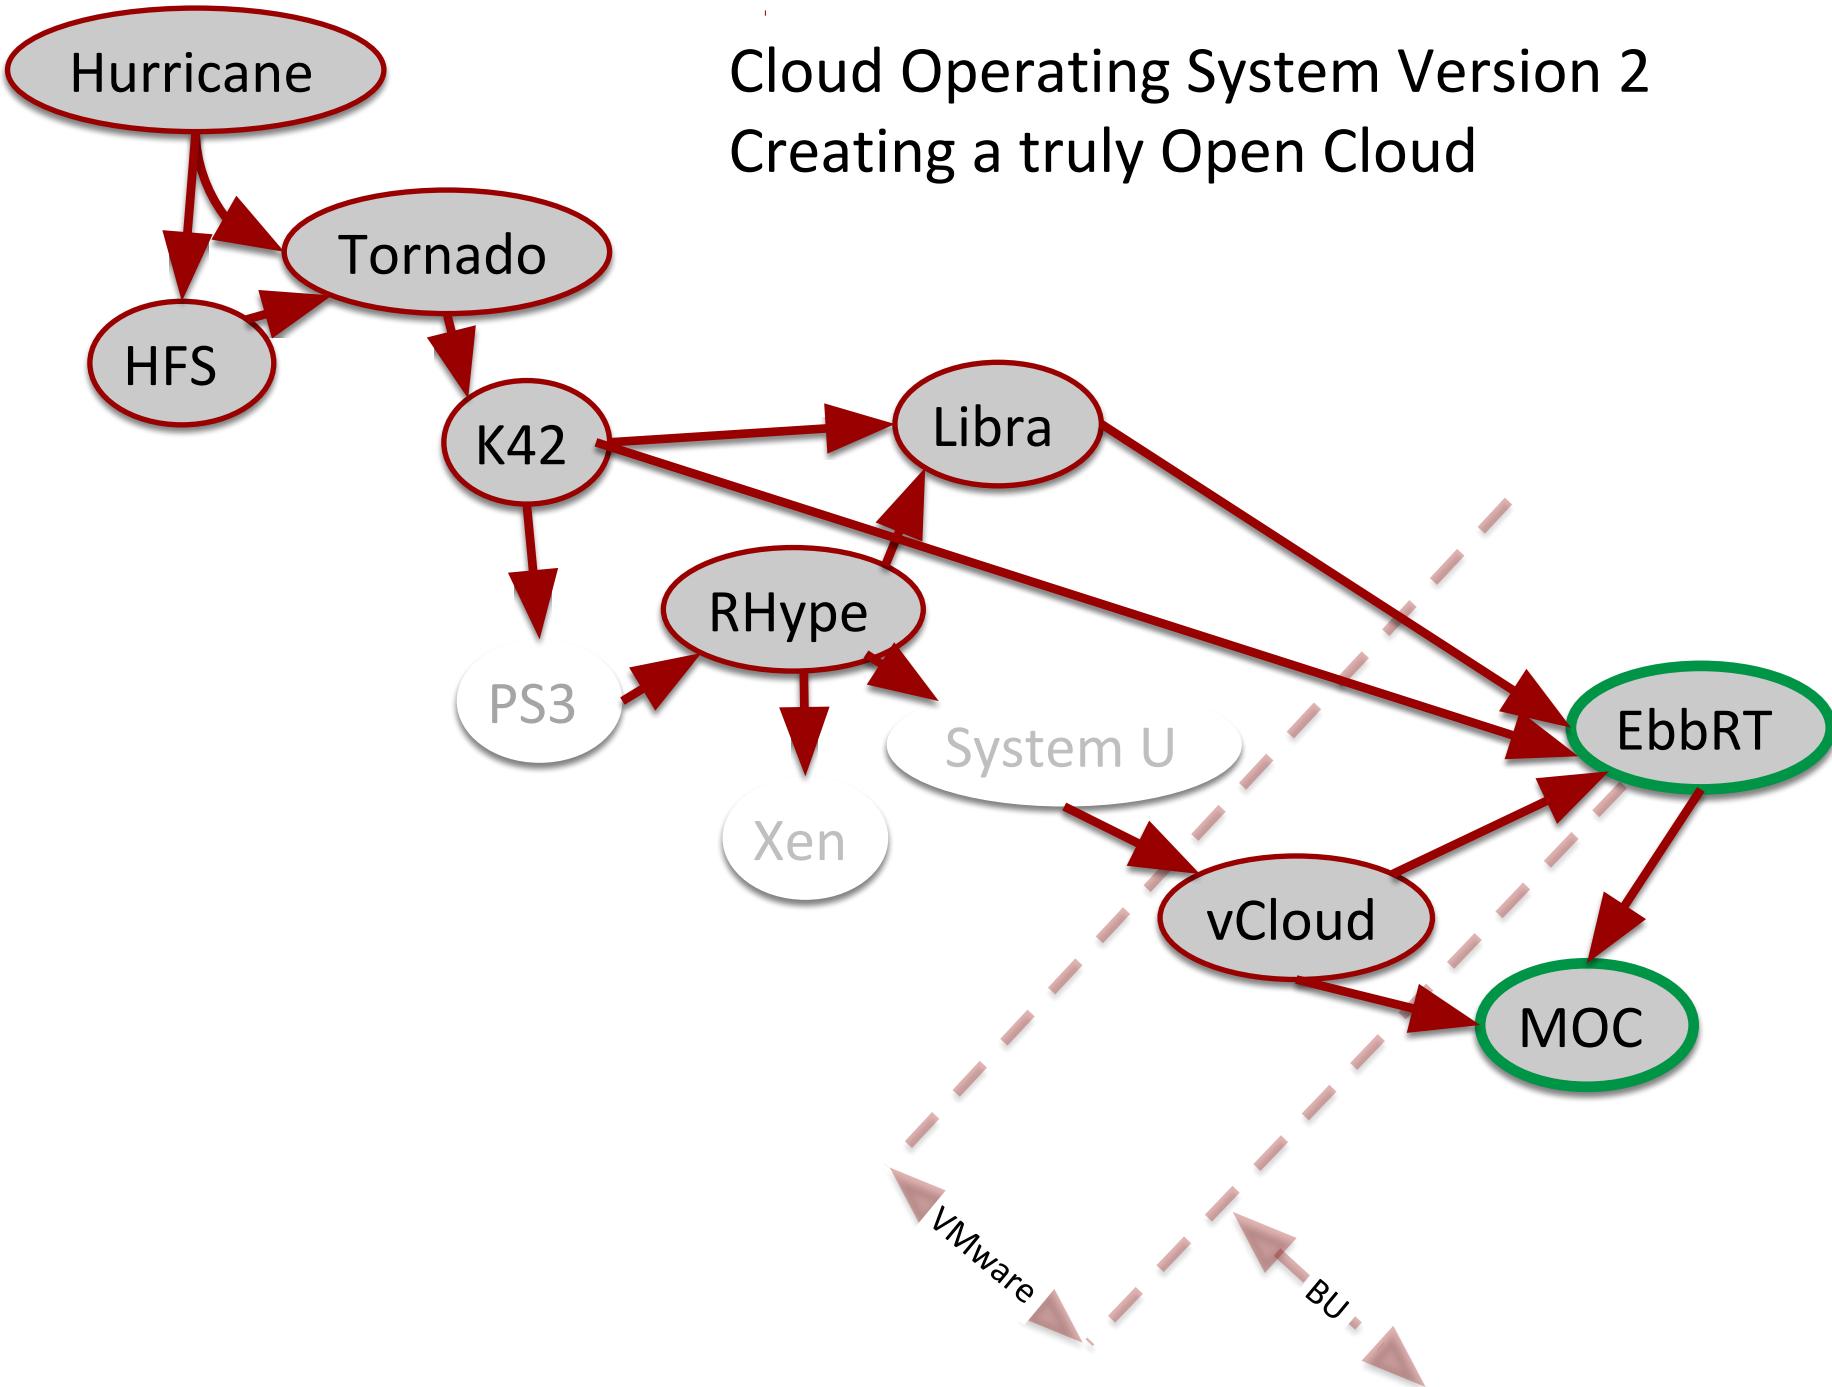


Computing is going to move to the cloud;



# Cloud Operating System Version 2

## Creating a truly Open Cloud



# Larry's background

- 40 years OS development at AT&T Bell Labs, Sony, Wang Labs, DEC, Red Hat
- Systems: AT&T System V Unix, Wang VSOS & Wang Unix, DEC VMS & Ultrix/Tru64 Unix, Red Hat Enterprise Linux(RHEL)
- Leading OS directions in the Red Hat Collaboratory@BU

# Larry's background(cont)

- Larry's primary kernel work has been:
  - memory management.
  - process control and scheduling.
  - interprocess communication and synchronization.
  - SMP and NUMA support.
  - crash dump support.
  - various file systems.

# Larry's background(cont)

- Open source kernel development
  - Years ago every company had their own OS
  - Today most OS development is done via the open source model.
    - large email lists to communicate design and implementation details.
    - Standardized design and coding styles.
    - designed by large committee of software and hardware engineers.
    - Lots of engineers contribute.

# What we are working on now

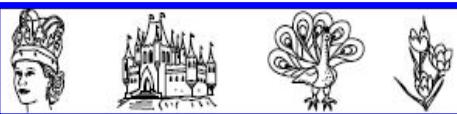
- Red Hat created Collaboratory@BU 4 years ago, we are working on:
  - Unikernel
  - Deterministic performance
  - OS security
  - AI control of OS policy
  - Partitioning Hypervisor
  - Automatic tuning of kernel parameters
- RH/IBM/BU Collaboration to create self-driven open platform from edge to cloud

# Outline

- What is an operating system
- How we got here
- Our background
- *What you will learn & course administration*

# Back to OS

Application



Operating System



Computer

## ***Multiplexing***

- creating an illusion of multiple (logical) resources from a single (physical) one

## ***Allocation***

- keep track of who has the right to use what

## ***Transformation***

- creating a new resource (logical) from existing resource (physical) primarily for “ease of use”

**An OS multiplexes, allocates, and transforms HW resources**

# Types of Multiplexing

## Time multiplexing

- time-sharing
- scheduling a serially-reusable resource among several users
  - e.g., CPU or printer

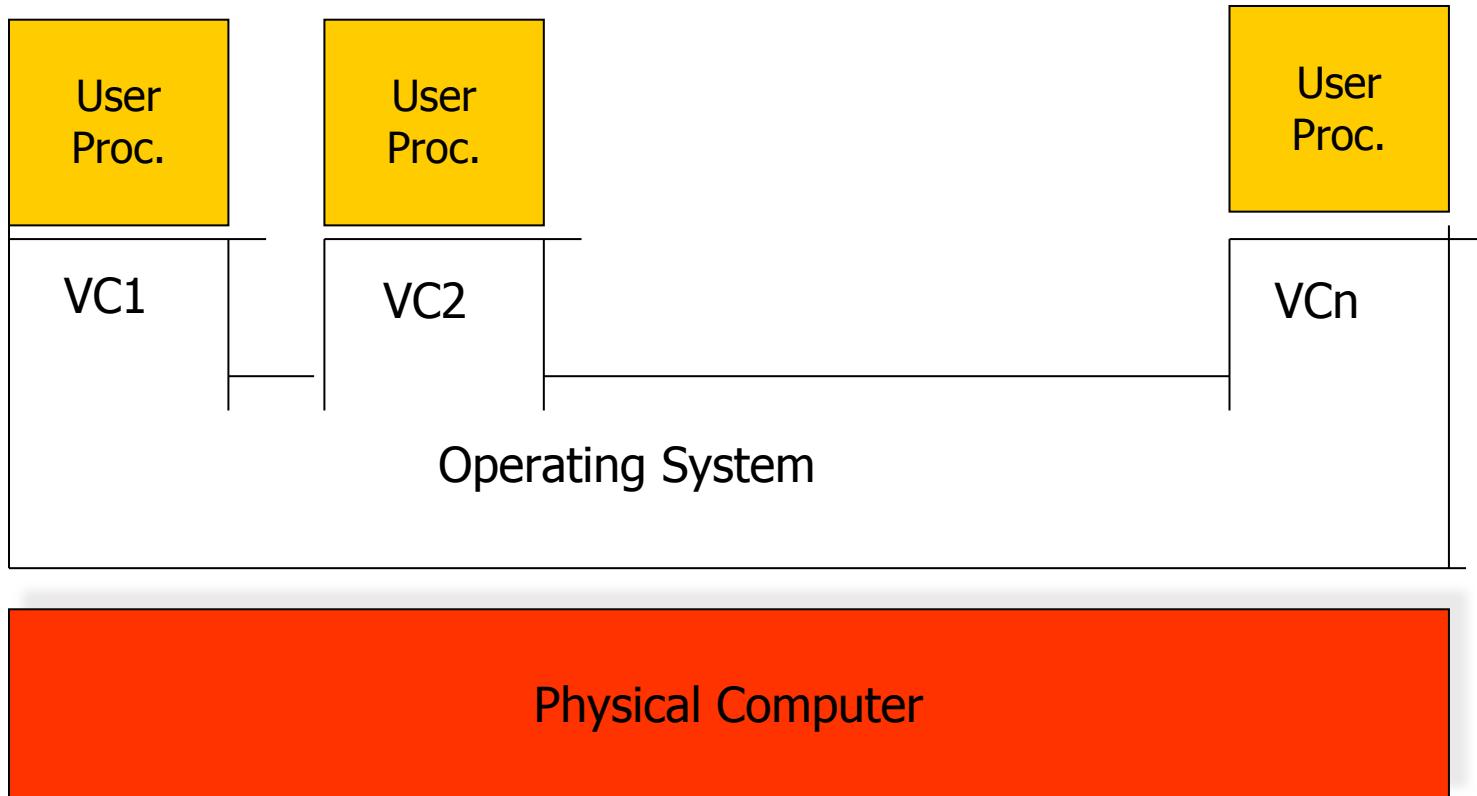
## Space multiplexing

- space-sharing
- dividing a multiple-use resource up among several users
  - e.g., memory, disk space

## We normally do both

- e.g., memory

# Virtual Computers



# OS Interface – Virtual Processors

**Nearly the same interface as the physical CPU**

**OS removes privileged operations**

- PSW determines if the code is either “user code” or “OS code”
- changes in status are strictly regulated...

**OS adds instructions (system calls)**

- create new virtual computers (processes)
- communicate with other VCs
- allocate memory
- perform input and output operations I/O
- access the file system

# OS Interface – Virtual Memory

- **Memory of the Virtual Computer is similar to the hardware memory (i.e., a sequence of words), and it is accessed the same way**
- **The OS divides up the memory into parts and gives each part to each virtual computer**
- **OS creates an illusion that each virtual computer has a memory starting from address 0x0000**
- **OS creates an illusion that the virtual computer has more memory than the physical memory**
- **The OS isolates and protects the memory between virtual computers**

# OS Interface – Virtual File System

- Secondary storage provides long-term storage for data
- Storage is done physically in term of disk sectors and virtually in terms of files
- The virtual computer sees a file system consisting of named files with arbitrary size

# OS Interface – Virtual I/O

- I/O operations of the virtual computer are completely different from the I/O operations of the physical computer
- The physical computer has devices with complex control and status registers
- In contrast, the virtual I/O is simple and easy to use
- In fact, in most OSes (e.g., UNIX) virtual I/O abstraction is almost identical to the file-system interface giving rise to uniformity of treatment with respect to all types of I/O devices including disks, terminals, printers, network connections
- Each VC sees a dedicated I/O device: the actual hardware is space/time multiplexed by the OS

# Key OS (e.g., Linux Kernel) services & interfaces

## Services

- System HW support
- CPU scheduling
- Memory management
- Device management
- Networking
- File systems
- Process environment

## Interfaces

- Configuration tools
- Processes & threads
- Virtual memory
- Device Drivers
- Sockets
- Directories and files
- System calls

# What you will learn

- How to write good system's code
- Key OS abstractions and interfaces
- A bit on the HW we are abstracting
- Processes/threads scheduling and synchronization
- Memory Management
- File systems
- Input and Output
- Virtualization and Cloud
- Security

# Projects

## Programming assignments

- 1) Shell (system calls)
- 2) Threads (parallel execution, scheduling)
- 3) Synchronization (semaphores, ...)
- 4) Memory (virtual memory, shared regions, ...)
- 5) File systems

We will also give detailed directions to download, build, boot and debug your own Linux kernel, and hope to provide a variant on 5, where you will build your own kernel file system.

- These are reputed to be really hard! Assume 20-30 hours of work for each.
- Don't CHEAT, you can talk about ideas, but no sharing code; you will need to fully explain your solution.

# EC 440 – Course Staff

- Orran Krieger & Larry Woodman
- Daniel Wilson (GTF)
- Jonathan Cameron & Ryan Sullivan (UTF)
- You - help each other, this course is tough... we will keep track of good answers on piazza

Should we have office hours?

# EC 440 Info

## Main resource: Piazza

- 1<sup>st</sup> stop for questions (check for duplicates!)
- Help each other with answers (**not solutions!**)
- Resources (lecture slides, etc.)
- Projects/homeworks/challenges will be posted there too
- Class Vidios will be posted on blackboard

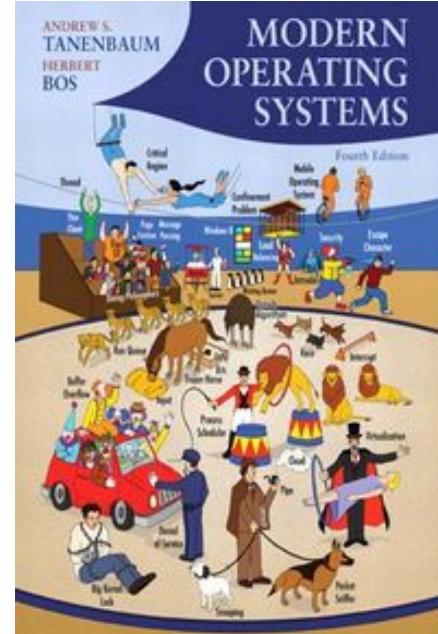
# Background

- This seems to be the “killer” course in CE
- We believe you should all be able to succeed... if you put in the work.
- Changing a few things from the past:
  - Handed out pre-challenge to help refresh you on programming in “c” and setting up your environment.
  - Developed a consistent environment with autograder
  - Next lecture Daniel will talk about programming
- We believe that you can do this, but we are not going to make it easier...

# Material

We will be informally using:

- Andrew S. Tanenbaum and Herbert Bos,  
Modern Operating Systems, Prentice-Hall,  
2015



# How we will mark

- Lots on the challenges ~75
- A fair bit on the exams ~20
- Some on the discussions on Piazza... ~5

We might add some bonus challenges, still thinking about it...